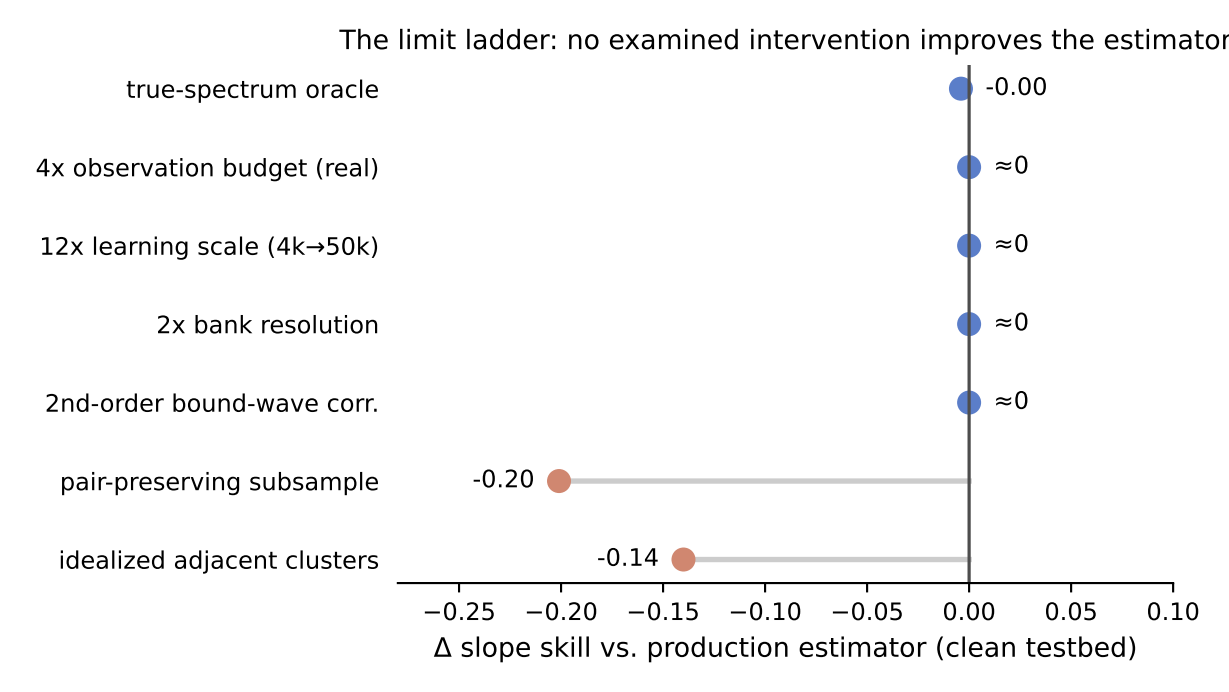
The limit ladder: no examined intervention improves the estimator
true-spectrum oracle
4x observation budget (real)
12x learning scale (4k→50k)
2x bank resolution
2nd-order bound-wave corr.
pair-preserving subsample
idealized adjacent clusters
-0.00
≈0
≈0
≈0
≈0
-0.20
-0.14
−0.25
−0.20
−0.15
−0.10
−0.05
0.00
0.05
0.10
Δ slope skill vs. production estimator (clean testbed)